RAG file-search: Azure Document Intelligence parsing + spreadsheet handling + parser-eval harness

Description

File-search ingestion used the stock RAG API loaders, so scanned/image PDFs weren't searchable and large spreadsheets timed out. This adds Azure AI Document Intelligence (OCR + layout) as the parser for pdf/docx/pptx/htm via a thin rag_api Docker overlay, with a graceful fallback to the native loaders when DI fails; routes native spreadsheets to a lean Excel loader that emits one self-describing record per row (with a large-workbook summary guard and adds a standalone parser-evaluation harness to benchmark DI vs the upstream loaders vs liteparse. Prompted by KARL UAT.

Environment

Normal text | Write

Type /ai for Atlassian Intelligence or @ to mention and notify someone.

Save Cancel

Acceptance Criteria

- [ ] Scanned/image PDFs in an agent's file-search are OCR'd and searchable (Azure DI).
- [ ] A DI failure on any file falls back to the native loader instead of erroring the upload.
- [ ] .xlsx is parsed natively into per-row records; oversized workbooks ingest fast via a summary (no timeout).
- [ ] parser-eval harness runs and produces a DI-vs-upstream-vs-liteparse comparison report.

Subtasks

Add subtask

Linked work items

Add linked work item

Activity

All | Comments | History | Work log | Smart Checklist History | Approvals



Add a comment...

Status update... Thanks... Agree...